

Text Representation Learning

Machine Learning Course - CS-433
Dec 6, 2023
Martin Jaggi
Last updated on: December 3, 2023

□


Motivation

Finding numerical representations for words is fundamental for all machine learning methods dealing with text data.

Goal: For each word, find mapping (embedding)

$$w_i \mapsto \mathbf{w}_i \in \mathbb{R}^K$$

Representation should capture semantics of the word.



Constructing good feature representations (= representation learning) benefits all ML applications.

The Co-Occurrence Matrix

A big corpus of un-labeled text can be represented as the co-occurrence counts
 $n_{ij} := \# \text{contexts where word } w_i \text{ occurs together with word } w_j.$

	1	1		
		3		
	1			
	2		1	
1				1
		1		
	1		1	1

Needs definition of

- Context e.g. document, paragraph, sentence, window
- Vocabulary

$$\mathcal{V} := \{w_1, \dots, w_D\}$$

For words $w_d = 1, 2, \dots, D$ and context words $w_n = 1, 2, \dots, N$, the co-occurrence counts n_{ij} form a very sparse $D \times N$ matrix.

Learning Word-Representations (Using Matrix Factorization)

Find a factorization of the cooccurrence matrix!

Typically uses log of the actual counts, i.e. $x_{dn} := \log(n_{dn})$.

We will aim to find \mathbf{W}, \mathbf{Z} s.t.

$$\mathbf{X} \approx \mathbf{W}\mathbf{Z}^\top$$

So for each pair of words (w_d, w_n) , we try to 'explain' their co-occurrence count by a numerical representation of the two words

- in fact by the inner product of the two feature vectors $\mathbf{W}_{d:}, \mathbf{Z}_{n:}$.

$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \frac{1}{2} \sum_{(d,n) \in \Omega} f_{dn} [x_{dn} - (\mathbf{W}\mathbf{Z}^\top)_{dn}]^2$$

where $\mathbf{W} \in \mathbb{R}^{D \times K}$ and $\mathbf{Z} \in \mathbb{R}^{N \times K}$ are tall matrices, having only $K \ll D, N$ columns.

The set $\Omega \subseteq [D] \times [N]$ collects the indices of non-zeros of the count matrix \mathbf{X} .

Each row of those matrices forms a representation of a word (\mathbf{W}) or a context word (\mathbf{Z}) respectively.


GloVe

This model is called GloVe, and is a variant of word2vec.

Weights f_{dn} : Give "importance" of each entry. Choosing $f_{dn} := 1$ is ok.

GloVe weight function:


$$f_{dn} := \min \{1, (n_{dn}/n_{\max})^\alpha\}, \quad \alpha \in [0; 1] \quad \text{e.g. } \alpha = \frac{3}{4}$$




Choosing K


K e.g. 50, 100, 500

Word Analogies





Verb tense



Country-Capital

Training

- Stochastic Gradient Descent (SGD)
- Alternating Least-Squares (ALS)

Open questions:

- Parallel and distributed training
- Does regularization help?

Alternative: Skip-Gram Model

(Original word2vec)


Uses binary classification (logistic regression objective), to separate real word pairs (w_d, w_n) from fake word pairs. Same inner product score = matrix factorization.

Given w_d , a context word w_n is

- real = appearing together in a context window of size 5
- fake = any word $w_{n'}$ sampled randomly: Negative sampling (also: Noise Contrastive Estimation)

Learning Representations of Sentences & Documents

Supervised: For a supervised task (e.g. predicting the emotion of a tweet), we can use matrixfactorization (below) or convolutional neural networks (see next weeks).



→ SemEval competition for tweet classification.

Unsupervised:

- Adding or averaging (fixed, given) word vectors
- Training word vectors such that adding/averaging works well
- Direct unsupervised training for sentences (appearing together with context sentences) instead of words

Fast Text

Matrix factorization to learn document/sentence representations (supervised).

Given a sentence $s_n = (w_1, w_2, \dots, w_m)$, let $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ be the bag-of-words representation of the sentence.

$$\min_{\mathbf{W}, \mathbf{Z}} \mathcal{L}(\mathbf{W}, \mathbf{Z}) := \sum_{s_n \text{ a sentence}} f(y_n \mathbf{W} \mathbf{Z}^T \mathbf{x}_n)$$

where $\mathbf{W} \in \mathbb{R}^{1 \times K}$, $\mathbf{Z} \in \mathbb{R}^{|\mathcal{V}| \times K}$ are the variables, and the vector $\mathbf{x}_n \in \mathbb{R}^{|\mathcal{V}|}$ represents our n -th training sentence.

Here f is a linear classifier loss function, and $y_n \in \{\pm 1\}$ is the classification label for sentence \mathbf{x}_n .

Language Models

Selfsupervised training:

Can a model generate text? - train classifier to predict the continuation (next word) of given text

- Multi-class:

Use soft-max loss function with a large number of classes $D = \text{vocabulary size}$

- Binary classification:

Predict if next word is real or fake (i.e. as in word2vec)


Impressive recent progress using large models, such as transformers

(e.g. GPT-2, GPT-3, chatGPT)

<https://transformer.huggingface.co/doc/gpt2-large>,

<https://chat.openai.com/>)

Arithmetic:



Reasoning:
Maksym Andriushchenko @maksym_andr
I got curious about this and tested ChatGPT on last year's exam from our
ML course at EPFL
(github.com/epfl/ml_course). Chain-of-thought evaluation with a ma-
jority vote over 5 trials gives 10/20
link: chatGPT on ML course exam

Further Pointers

1. word2vec:

code: code.google.com/p/word2vec/

paper:

"Distributed representations of words and phrases and their compositional-
ity" - T Mikolov, I Sutskever, K Chen, GS Corrado, J Dean. NIPS 2013

2. GloVe:

code and vectors: nlp.stanford.edu/projects/glove/

paper:

"GloVe: Global Vectors for Word Representation" - Pennington, J., Socher,
R., Manning, C. D.. EMNLP 2014

3. FastText & sent2vec

code: github.com/facebookresearch/fastText

papers:

"Bag of Tricks for Efficient Text Classification" - Joulin, A., Grave, E.,
Bojanowski, P., Mikolov, T. - EC-ACL, 2017.

"Enriching Word Vectors with Subword Information" - Bojanowski,
P., Grave, E., Joulin, A., Mikolov, T. - TACL, 2017.

"Unsupervised Learning of Sentence Embeddings using Compositional n-
Gram Features" - Pagliardini, M., Gupta, P., Jaggi, M. NAACL 2018.

4. Write with transformers:

code and demo: transformer.huggingface.co/doc/gpt2-large

5. ChatGPT

demo: chat.openai.com/